



#ASLI ENGINEERING

Canary Deployments



BY

ARPIT BHAYANI

Canary Deployments

Canary deployment is a pattern that allows us to rollout code, features, changes to an initial subset of users/servers before we take it to 100%.

roll out to canary



if all good, roll out to 100% of servers



if something is wrong, rollback

Canary deployment thus acts as an Early Warning Indicator so that we could avoid a massive outage

The flow would be:

1. deploy the new version to one server
2. monitor the vitals: CPU, RAM, error rates
3. test the changes explicitly on canary
4. if all good, rollout to remaining servers
5. if something is wrong, rollback

Why the name "canary deployment" ?

In 1920s during coal mining, the miners used to carry caged **canaries** into the coal mines.




if there was a high level of toxic gas, the canary would **die**, alerting the miners to evacuate.

In our canary deployment the few servers/users we roll out to our canaries for us that

- provides early warnings if something goes wrong


How canary deployments are implemented ?

We create a small parallel infrastructure and put a proxy like load balancer or API gateway in front



Pros of doing a canary deployment


- allows us to test the changes in production with real users



The fleet can be
logically separate
or physically
separate

- rollbacks are much faster

Because the new version of code is deployed on only a few machines, rolling back our changes requires us to rollback only those few machines



- minimal blast radius

if the new version of code has "bugs"
with canary deployment it would affect a
very limited set of requests

- zero downtime deployments

incremental rollout: 1%, 5%, 10%, 20%, 50%, 100%.

- we could deploy even when we are unsure about the new release

- we can use canary setup to do A/B testing

eg: search v1 and search v2

monitor CTR side by side

Note: Selection of servers / users can be more sophisticated

- geographical

- sticky + random selection

- user cohorts

- internal employees

- random selection

Selection criteria is usecase specific

- beta users

Cons of doing a canary deployment

- engineers will be **habituated** to test in **production**

Testing should never happen in production, but if we have a canary deployment we are **encouraging** engineers to test in the production environment, and slowly we may find most testing happening in production.


- **architecting** a canary deployment is **complex**

Extra infra components: LB, API Gateway, separate scaling policy, instance/container launch config etc.

- **parallel monitoring** setup


Observability is super.important in canary setup

Side by side comparison of vitals - CPU, RAM, response times, exceptions, etc.



You absolutely need canary ↓

When you re-wrote your service



You would not do a 0 to 1 switch from one service to another.

This has to be done incrementally

100% of traffic will be moved to new Auth Service written in Golang only after you have 100% confidence about the correctness of implementation