



**#ASLI ENGINEERING**

# GitHub Outage: Downtime due to reversing an Index

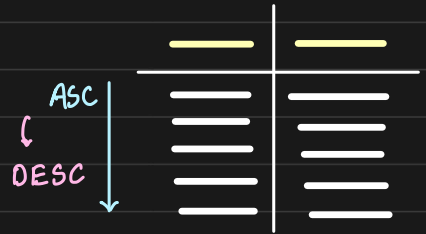


**BY**

**ARPIT BHAYANI**

### Downtime due to "reversing" an index

GitHub ran a database migration  
to flip the order of an index  
to improve query performance



commits		
id	user_id	date

Get commits on a repository  
in descending order

Get commits ordered by date and for each day order by user

↓                      ↓

**DESC**                      **ASC**

# ARPIT BHAYANI

Hence to answer the query efficiently

We would need to have an index which is  
ordered by **Date** in **DESC** order and **UserID** in **ASC** order

→ Supported by  
MySQL 8

Hence, we have to "flip" the index

Reversing the index caused **Full Table Scan**

for a query generated by their ORM (Active Record)

Load on  
the DB



1. Creating a new index itself would require a **Full Table Scan**
2. Say, you changed the order of the index, but

1. some other query needed **DATE, UserID ASC**

which is much more frequent would now become inefficient

2. some queries might not use this new index

because DB engine thought it would be "inefficient"

↳ because of the query we provided

↳ on its own estimations

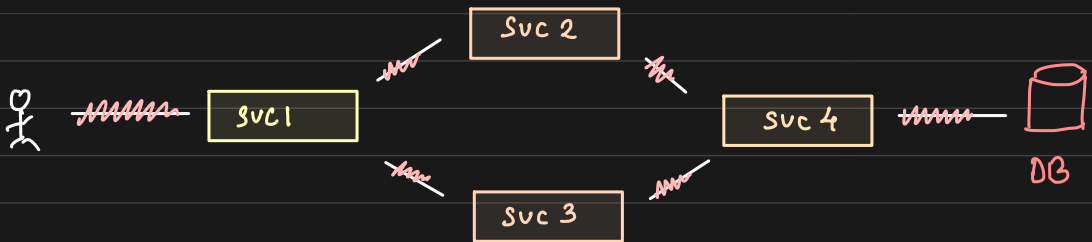
We can use **INDEX HINTS** to tell our DB  
to use a particular index for a particular  
query

```
SELECT .....  
FROM .....  
USE INDEX (idx_.....);
```

## Cascading Effect

Because a query was doing a Full Table Scan

- ↳ Put load on the DB
- ↳ Response time of the service increased
- ↳ lead to timeouts of the request
- ↳ Response time of depending services increased
- ↳ timeouts cascaded !!



Note: Whenever there is a synchronous dependency b/w services  
there is always a chance of Cascading Failures

## Key Takeaways:

1. Do not blindly trust your ORM ↗ Active Record, Django ORM, SQLAlchemy, Hibernate

ORM makes our lives simpler, but they may not generate the most optimized SQL queries

Periodically audit the queries generated by ORMs

2. Always check the Query Execution Plan

Run your migration on staging / QA environment

and see the Query execution plan for related queries

Check for any aberration or deviation

3. Audit the queries and indexes they use

prepare an inventory of queries and indexes they use

so that we can quickly test for regressions